


✓ Importing the dataset

```
#Install kaggle
!pip install -q kaggle
```

```
from google.colab import files
files.upload()
```

 no files selected Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.
Saving kaggle.json to kaggle (1).json
f'kaggle (1).json'.


```
#Create a kaggle folder
! mkdir ~/.kaggle
```

 mkdir: cannot create directory '/root/.kaggle': File exists


```
#Copy the kaggle.json to created folder
! cp kaggle.json ~/.kaggle/
```

```
#Permission for the json to act
! chmod 600 ~/.kaggle/kaggle.json
```

```
#Download Audio MNIST dataset
!kaggle datasets download sripaadsrinivasan/audio-mnist
```

 Dataset URL: <https://www.kaggle.com/datasets/sripaadsrinivasan/audio-mnist>
License(s): CC0-1.0
audio-mnist.zip: Skipping, found more recently modified local copy (use --force)

```
#Unzip to access audio files
!unzip audio-mnist.zip
```

 inflating: data/51/1_51_32.wav
inflating: data/51/1_51_33.wav
inflating: data/51/1_51_34.wav

```
inflating: data/51/1_51_34.wav
inflating: data/51/1_51_35.wav
inflating: data/51/1_51_36.wav
inflating: data/51/1_51_37.wav
inflating: data/51/1_51_38.wav
inflating: data/51/1_51_39.wav
inflating: data/51/1_51_4.wav
inflating: data/51/1_51_40.wav
inflating: data/51/1_51_41.wav
inflating: data/51/1_51_42.wav
inflating: data/51/1_51_43.wav
inflating: data/51/1_51_44.wav
inflating: data/51/1_51_45.wav
inflating: data/51/1_51_46.wav
inflating: data/51/1_51_47.wav
inflating: data/51/1_51_48.wav
inflating: data/51/1_51_49.wav
inflating: data/51/1_51_5.wav
inflating: data/51/1_51_6.wav
inflating: data/51/1_51_7.wav
inflating: data/51/1_51_8.wav
inflating: data/51/1_51_9.wav
inflating: data/51/2_51_0.wav
inflating: data/51/2_51_1.wav
inflating: data/51/2_51_10.wav
inflating: data/51/2_51_11.wav
inflating: data/51/2_51_12.wav
inflating: data/51/2_51_13.wav
inflating: data/51/2_51_14.wav
inflating: data/51/2_51_15.wav
inflating: data/51/2_51_16.wav
inflating: data/51/2_51_17.wav
inflating: data/51/2_51_18.wav
inflating: data/51/2_51_19.wav
inflating: data/51/2_51_2.wav
inflating: data/51/2_51_20.wav
inflating: data/51/2_51_21.wav
inflating: data/51/2_51_22.wav
inflating: data/51/2_51_23.wav
inflating: data/51/2_51_24.wav
inflating: data/51/2_51_25.wav
inflating: data/51/2_51_26.wav
inflating: data/51/2_51_27.wav
inflating: data/51/2_51_28.wav
inflating: data/51/2_51_29.wav
inflating: data/51/2_51_3.wav
inflating: data/51/2_51_30.wav
inflating: data/51/2_51_31.wav
inflating: data/51/2_51_32.wav
inflating: data/51/2_51_33.wav
inflating: data/51/2_51_34.wav
```

```

inflating: data/51/2_51_35.wav
inflating: data/51/2_51_36.wav
inflating: data/51/2_51_37.wav
inflating: data/51/2_51_38.wav
inflating: data/51/2_51_39.wav
inflating: data/51/2_51_4.wav
inflating: data/51/2_51_40.wav

```

✓ Play and visualize the chosen audio file

```

import numpy as np
import librosa
import librosa.display as dsp
import matplotlib.pyplot as plt
from IPython.display import Audio, display, clear_output
import os
import ipywidgets as widgets

class SequentialAudioPlayer:
    def __init__(self, base_path="data"):
        #Initialize player with file paths and UI controls
        self.base_path = base_path
        self.current_speaker = None
        self.current_digit = 0
        #Define navigation order for recordings
        self.recording_order = ['0','1','10','11','12','13','14','15','16','17','18','19','20','21','22','23','24','25','26','27','28','29','30','31','32','33','34','35','36','37','38','39','4','40','41','42','43','44','45','46','47','48','49','5','6','7','8','9']

        self.current_index = 0
        self.total_recordings = 50 #Recordings per speaker

        #Initialize UI input widgets and controls
        self.speaker_input = widgets.IntText(
            value = 1,
            description = 'Speaker:',
            min = 1,
            max = 60
        )

        self.digit_input = widgets.IntText(
            value = 0,
            description = 'Digit:',

```

```

        min = 0,
        max = 9
    )

    self.play_button = widgets.Button(description="Play")
    self.next_button = widgets.Button(description="Next")
    self.prev_button = widgets.Button(description="Previous")
    self.jump_input = widgets.IntText(
        value = 1,
        description = 'Jump to:',
        min = 1,
        max = 500
    )
    self.jump_button = widgets.Button(description = "Jump")

    #Labels to track position and the current file
    self.position_label = widgets.Label()
    self.filename_label = widgets.Label()

    #Connect button click handlers
    self.play_button.on_click(self.play_current)
    self.next_button.on_click(self.next_audio)
    self.prev_button.on_click(self.prev_audio)
    self.speaker_input.observe(self.speaker_changed, names = 'value')
    self.digit_input.observe(self.digit_changed, names = 'value')
    self.jump_button.on_click(self.jump_to_recording)

    #Building the player
    self.controls = widgets.VBox([
        widgets.HBox([
            self.speaker_input,
            self.digit_input,
            self.prev_button,
            self.play_button,
            self.next_button
        ]),
        widgets.HBox([
            self.jump_input,
            self.jump_button
        ]),
        self.position_label,
        self.filename_label
    ])

    def get_current_recording_number(self):

```

```
#Convert position to recording number (1-500)
return self.current_digit * 50 + self.current_index + 1

def get_file_path(self):
    #File path from current position
    speaker = self.current_speaker
    digit = self.current_digit
    index = self.recording_order[self.current_index]
    speaker_str = f"0{speaker}" if speaker < 10 else str(speaker)
    filename = f"{digit}_{speaker_str}_{index}.wav"
    return os.path.join(self.base_path, speaker_str, filename), filename

def update_labels(self):
    #Current position label
    file_path, filename = self.get_file_path()
    current_recording = self.get_current_recording_number()
    self.position_label.value = f"Recording {current_recording}/500"
    self.filename_label.value = f"Current file: {filename}"

def speaker_changed(self, change):
    #Reset position when speaker changes
    self.current_speaker = change.new
    self.current_digit = self.digit_input.value
    self.current_index = 0
    self.update_labels()
    self.display_current()

def digit_changed(self, change):
    #Reset position when digit changes
    self.current_digit = change.new
    self.current_index = 0
    self.update_labels()
    self.display_current()

def jump_to_recording(self, b=None):
    #Navigate to chosen recording number
    target = self.jump_input.value
    if 1 <= target <= 500:
        self.current_digit = (target - 1) // 50
        self.current_index = (target - 1) % 50
        self.digit_input.value = self.current_digit
        self.update_labels()
        self.display_current()

def play_current(self, b=None):
```

```
#Play current audio file
file_path, _ = self.get_file_path()
if os.path.exists(file_path):
    data, sr = librosa.load(file_path)
    display(Audio(data=data, rate=sr, autoplay=True))

def next_audio(self, b=None):
    #Next recording
    self.current_index += 1
    if self.current_index >= 50:
        self.current_index = 0
        self.current_digit += 1
        if self.current_digit > 9:
            self.current_digit = 0
            new_speaker = self.current_speaker + 1
            if new_speaker <= 60:
                self.speaker_input.value = new_speaker
    self.digit_input.value = self.current_digit
    self.update_labels()
    self.display_current()

def prev_audio(self, b=None):
    #Previous recording
    self.current_index -= 1
    if self.current_index < 0:
        self.current_index = 49
        self.current_digit -= 1
        if self.current_digit < 0:
            self.current_digit = 9
            new_speaker = self.current_speaker - 1
            if new_speaker >= 1:
                self.speaker_input.value = new_speaker
    self.digit_input.value = self.current_digit
    self.update_labels()
    self.display_current()




def display_current(self):
    #Visualize audio
    clear_output(wait=True)
    display(self.controls)


    file_path, _ = self.get_file_path()
    if os.path.exists(file_path):
        data, sr = librosa.load(file_path)
        plt.figure(figsize=(12, 4))
```

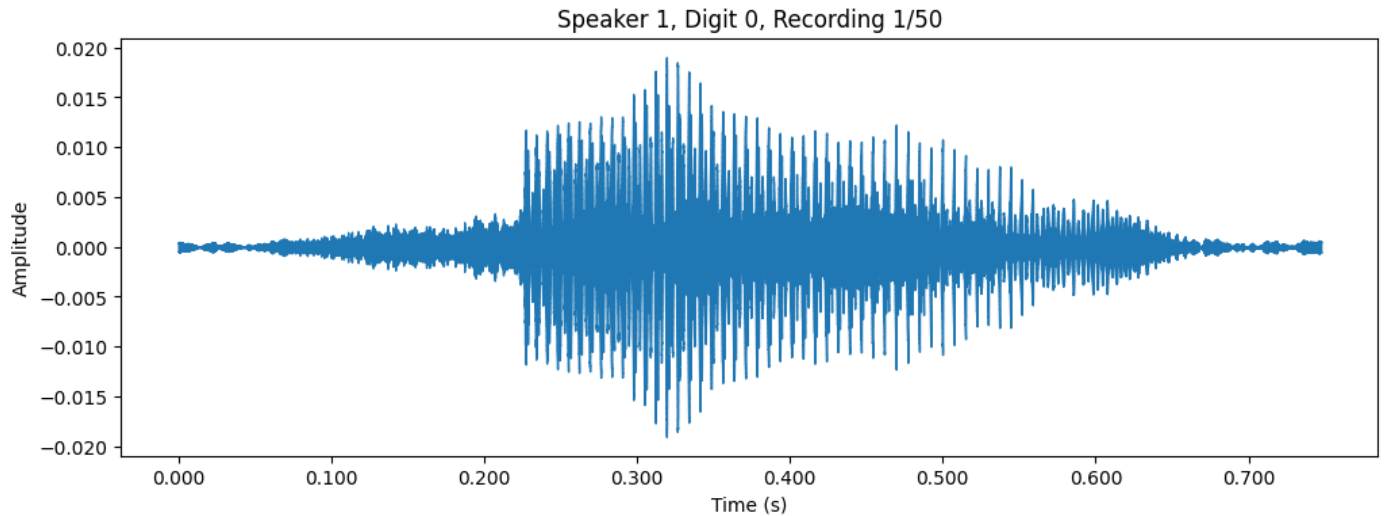
```
        dsp.waveshow(data, sr=sr)
        plt.title(f"Speaker {self.current_speaker}, Digit {self.current_digit}")
        plt.xlabel("Time (s)")
        plt.ylabel("Amplitude")
        plt.show()
    else:
        print(f"File not found: {file_path}")

def create_player():
    #Display interface
    player = SequentialAudioPlayer()
    player.current_speaker = player.speaker_input.value
    player.display_current()
    return player

if __name__ == "__main__":
    player = create_player()
```


 Speaker: 
 Digit: 

Jump to: 
 Jump



✓ Extracting MFCC Coefficients from the audio files

```

import numpy as np
import librosa
import librosa.display
import matplotlib.pyplot as plt
%matplotlib inline
import os
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
import torch
from torch.utils.data import Dataset, DataLoader
import IPython.display as ipd

```



```
def extract_mfcc(audio_path, sample_rate=16000, window_length=0.03, hop_length=0.01):
    #Extract MFCC features from audio file
    try:
        audio, sr = librosa.load(audio_path, sr=sample_rate)
        #Calculate FFT parameters
        n_fft = int(window_length * sample_rate)
        hop_length_samples = int(hop_length * sample_rate)

        #Compute MFCC features
        mfcc = librosa.feature.mfcc(
            y = audio,
            sr = sample_rate,
            n_mfcc = n_mfcc,
            n_fft = n_fft,
            hop_length = hop_length_samples
        )
        return mfcc.T
    except Exception as e:
        print(f"Error processing {audio_path}: {str(e)}")
        return None

def visualize_mfcc(audio_path, sample_rate=16000):
    #Visualize audio and extracted MFCC Features
    plt.clf()

    #Load and process audio
    audio, sr = librosa.load(audio_path, sr=sample_rate)
    mfcc = extract_mfcc(audio_path, sample_rate)

    fig, (ax1, ax2) = plt.subplots(2, 1, figsize=(12, 8))


    #Plot waveform and MFCC
    librosa.display.waveshow(y=audio, sr=sr, ax=ax1)
    ax1.set_title('Waveform')
    img = librosa.display.specshow(mfcc.T, x_axis='time', ax=ax2)
    fig.colorbar(img, ax=ax2, format='%+2.0f dB')
    ax2.set_title('MFCC')

    plt.tight_layout()
    plt.show()
    return ipd.Audio(audio, rate=sr)

#Test visualization on sample file
sample_file = "data/01/0_01_0.wav"
```

```
visualize_mfcc(sample_file)
```

 <Figure size 640x480 with 0 Axes>



-0:00

✓ Prepare the dataset

```
#Import required libraries
import numpy as np
import librosa
import os
from sklearn.model_selection import train_test_split
from torch.utils.data import Dataset, DataLoader
from concurrent.futures import ThreadPoolExecutor
from tqdm import tqdm

def process_file(args):
    #Pad/truncate MFCC features to fixed length
    file_path, speaker = args
    mfcc = extract_mfcc(file_path)
    if mfcc is not None:
        #Standardize MFCC length to 40 frames (explained why in report)
        if mfcc.shape[0] < 40:
            pad_width = ((0, 40 - mfcc.shape[0]), (0, 0))
            mfcc = np.pad(mfcc, pad_width, mode='constant')
        else:
            mfcc = mfcc[:40, :]
        return mfcc, speaker - 1
    return None

def prepare_dataset(base_path="data", test_size=0.2, verbose=True):
    if verbose:
        print("Preparing dataset...")

    #Collect all audio files
    file_list = []
    for speaker in range(1, 61):
        speaker_dir = f"{base_path}/{str(speaker) if speaker < 10 else str(speaker)}"
        if not os.path.exists(speaker_dir):
            continue
        for file in os.listdir(speaker_dir):
            if file.endswith('.wav'):
                file_list.append((os.path.join(speaker_dir, file), speaker))

    #Extract features
    features = []
    labels = []
```

```

with ThreadPoolExecutor(max_workers=8) as executor:
    results = list(tqdm(executor.map(process_file, file_list),
                          total=len(file_list),
                          desc="Processing audio files"))

#Collect successful results
for result in results:
    if result is not None:
        features.append(result[0])
        labels.append(result[1])

#Convert to numpy arrays and split dataset
X = np.array(features)
y = np.array(labels)
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=test_size, random_state=42, stratify=y
)

#Normalize features
for i in range(X_train.shape[2]):
    mean = X_train[:, :, i].mean()
    std = X_train[:, :, i].std()
    X_train[:, :, i] = (X_train[:, :, i] - mean) / std
    X_test[:, :, i] = (X_test[:, :, i] - mean) / std

if verbose:
    print("\nDataset Info:")
    print(f"Total samples: {len(features)}")
    print(f"Training samples: {len(X_train)}")
    print(f"Test samples: {len(X_test)}")
    print(f"Feature dimensions: {X_train.shape}")
    print(f"Number of speakers: {len(np.unique(y_train))}")
    print(f"Files processed: {len(results)}")

return X_train, X_test, y_train, y_test

class SpeakerDataset(Dataset):
    #PyTorch dataset for speaker recognition
    def __init__(self, features, labels):
        self.features = torch.FloatTensor(features)
        self.labels = torch.LongTensor(labels)

    def __len__(self):
        return len(self.labels)

```

```

def __getitem__(self, idx):
    return self.features[idx], self.labels[idx]

#Create and load datasets
if __name__ == "__main__":
    X_train, X_test, y_train, y_test = prepare_dataset()
    train_dataset = SpeakerDataset(X_train, y_train)
    test_dataset = SpeakerDataset(X_test, y_test)
    train_loader = DataLoader(train_dataset, batch_size=32, shuffle=True)
    test_loader = DataLoader(test_dataset, batch_size=32, shuffle=False)

```

➡ Preparing dataset...

Processing audio files: 100%|██████████| 30000/30000 [03:53<00:00, 128.52it/s]

Dataset Info:
 Total samples: 30000
 Training samples: 24000
 Test samples: 6000
 Feature dimensions: (24000, 40, 13)
 Number of speakers: 60
 Files processed: 30000

✓ Creating and Training the CNN Model

```

#Import PyTorch modules
import torch
import torch.nn as nn
from torch.optim import Adam
from torch.optim.lr_scheduler import ReduceLROnPlateau
from tqdm import tqdm

class SpeakerCNN(nn.Module):
    def __init__(self):
        super().__init__()
        #CNN architecture => 3 convolutional blocks
        self.conv = nn.Sequential(
            #First convolutional block: 32 filters
            nn.Conv2d(1, 32, 3, padding=1),
            nn.BatchNorm2d(32),
            nn.ReLU(),
            nn.MaxPool2d(2),

            #Second

```

```

        nn.Conv2d(32, 64, 3, padding=1),
        nn.BatchNorm2d(64),
        nn.ReLU(),
        nn.MaxPool2d(2),

        #Third
        nn.Conv2d(64, 128, 3, padding=1),
        nn.BatchNorm2d(128),
        nn.ReLU(),
        nn.MaxPool2d(2)
    )

    #Classifier => 2-layer fully connected with dropout
    self.fc = nn.Sequential(
        nn.Linear(128 * 5 * 1, 256),
        nn.ReLU(),
        nn.Dropout(0.5),
        nn.Linear(256, 60)
    )

    def forward(self, x):
        x = x.unsqueeze(1) #Add channel dimension
        x = self.conv(x)
        x = x.view(-1, 128 * 5 * 1)
        return self.fc(x)

def train(model, train_loader, test_loader, epochs=20):
    #Setup training
    device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
    print(f"Using device: {device}")
    model = model.to(device)
    criterion = nn.CrossEntropyLoss()
    optimizer = Adam(model.parameters(), lr=0.001)
    scheduler = ReduceLROnPlateau(optimizer, mode='min', patience=3, factor=0.5)
    best_accuracy = 0
    best_model_state = None

    #Training
    for epoch in range(epochs):
        model.train()
        train_loss = 0
        for features, labels in tqdm(train_loader, desc=f'Epoch {epoch+1}/{epochs}'):
            features, labels = features.to(device), labels.to(device)
            optimizer.zero_grad()
            outputs = model(features)
            _ = criterion(outputs, labels)
            train_loss += _

```

```
#Initialize and train model
model = SpeakerCNN()
trained model = train(model, train_loader, test_loader)
```

```
Using device: cpu
Epoch 1/20: 100%|██████████| 750/750 [00:40<00:00, 18.65it/s]

Epoch 1 - Average Loss: 2.4339
Validation Accuracy: 68.67%
New best accuracy: 68.67%
Epoch 2/20: 100%|██████████| 750/750 [00:41<00:00, 18.28it/s]
```


Epoch 2 – Average Loss: 0.9985
Validation Accuracy: 85.68%
New best accuracy: 85.68%
Epoch 3/20: 100%|██████████| 750/750 [00:41<00:00, 17.96it/s]

Epoch 3 – Average Loss: 0.6550
Validation Accuracy: 84.67%
Epoch 4/20: 100%|██████████| 750/750 [00:40<00:00, 18.45it/s]

Epoch 4 – Average Loss: 0.4988
Validation Accuracy: 92.40%
New best accuracy: 92.40%
Epoch 5/20: 100%|██████████| 750/750 [00:40<00:00, 18.49it/s]

Epoch 5 – Average Loss: 0.4120
Validation Accuracy: 93.77%
New best accuracy: 93.77%
Epoch 6/20: 100%|██████████| 750/750 [00:41<00:00, 18.23it/s]

Epoch 6 – Average Loss: 0.3475
Validation Accuracy: 96.10%
New best accuracy: 96.10%
Epoch 7/20: 100%|██████████| 750/750 [00:40<00:00, 18.36it/s]

Epoch 7 – Average Loss: 0.3131
Validation Accuracy: 95.00%
Epoch 8/20: 100%|██████████| 750/750 [00:40<00:00, 18.51it/s]

Epoch 8 – Average Loss: 0.2601
Validation Accuracy: 97.02%
New best accuracy: 97.02%
Epoch 9/20: 100%|██████████| 750/750 [00:39<00:00, 18.88it/s]

Epoch 9 – Average Loss: 0.2236
Validation Accuracy: 95.83%
Epoch 10/20: 100%|██████████| 750/750 [00:41<00:00, 18.18it/s]

Epoch 10 – Average Loss: 0.2037
Validation Accuracy: 96.38%
Epoch 11/20: 100%|██████████| 750/750 [00:40<00:00, 18.58it/s]

Epoch 11 – Average Loss: 0.1924
Validation Accuracy: 97.28%
New best accuracy: 97.28%
Epoch 12/20: 100%|██████████| 750/750 [00:40<00:00, 18.55it/s]

Epoch 12 – Average Loss: 0.1670
Validation Accuracy: 96.78%
Epoch 13/20: 100%|██████████| 750/750 [00:40<00:00, 18.39it/s]

Epoch 13 – Average Loss: 0.1654

```

Epoch 100, Average Loss: 0.1234,
Validation Accuracy: 0.789%

```

✓ Test the CNN Model

```

import torch
import numpy as np
from sklearn.metrics import precision_recall_fscore_support, classification_report
import random
import os

def test_model(model, test_loader, device):
    #Evaluate model performance on test set
    model.eval()
    all_predictions = []
    all_labels = []
    correct = total = 0

    with torch.no_grad():
        for features, labels in test_loader:
            features, labels = features.to(device), labels.to(device)
            outputs = model(features)
            predictions = outputs.argmax(1)

            correct += (predictions == labels).sum().item()
            total += len(labels)
            all_predictions.extend(predictions.cpu().numpy())
            all_labels.extend(labels.cpu().numpy())

    #Calculate and print performance metrics
    accuracy = 100 * correct / total
    precision, recall, f1, _ = precision_recall_fscore_support(
        all_labels, all_predictions, average = 'weighted'
    )

    print("\nModel Performance Metrics:")
    print(f"Test Accuracy: {accuracy:.2f}%")
    print(f"Weighted Precision: {precision:.4f}")
    print(f"Weighted Recall: {recall:.4f}")
    print(f"Weighted F1-Score: {f1:.4f}")
    print("\nDetailed Classification Report:")
    print(classification_report(all_labels, all_predictions))

    return all_predictions, all_labels

```

```

def predict_speaker(model, audio_file, device):
    #Predict speaker identity from audio file
    mfcc = extract_mfcc(audio_file)

    #Standardize input length
    if mfcc.shape[0] < 40:
        pad_width = ((0, 40 - mfcc.shape[0]), (0, 0))
        mfcc = np.pad(mfcc, pad_width, mode='constant')
    else:
        mfcc = mfcc[:40, :]

    #Normalize features (X-mean/s.d)
    mfcc = (mfcc - mfcc.mean(axis=0)) / (mfcc.std(axis=0) + 1e-8)
    mfcc_tensor = torch.FloatTensor(mfcc).unsqueeze(0).to(device)

    model.eval()
    with torch.no_grad():
        output = model(mfcc_tensor)
        temperature = 1.5 #Calibration parameter
        scaled_output = output / temperature
        probabilities = torch.nn.functional.softmax(scaled_output, dim=1)
        predicted_speaker = probabilities.argmax(1).item() + 1
        confidence = probabilities.max().item() * 100

    return predicted_speaker, confidence, probabilities.cpu().numpy()[0]

def get_audio_file_path(speaker_id, base_path="data"):
    #Get random audio sample from speaker
    speaker_str = f"{'0' + str(speaker_id) if speaker_id < 10 else str(speaker_id)}"
    speaker_dir = os.path.join(base_path, speaker_str)

    if os.path.exists(speaker_dir):
        audio_files = [f for f in os.listdir(speaker_dir) if f.endswith('.wav')]
        if audio_files:
            return os.path.join(speaker_dir, random.choice(audio_files))
    return None

def test_random_sample(model, test_loader, device, base_path="data"):
    #Test model on random sample
    batch_idx = random.randint(0, len(test_loader) - 1)
    for i, (features, labels) in enumerate(test_loader):
        if i == batch_idx:
            sample_idx = random.randint(0, len(features) - 1)
            test_feature = features[sample_idx:sample_idx+1]

```

```

true_speaker = labels[sample_idx].item() + 1
audio_file = get_audio_file_path(true_speaker, base_path)

#Predict
test_feature = test_feature.to(device)
model.eval()
with torch.no_grad():
    output = model(test_feature)
    predicted_speaker = output.argmax(1).item() + 1
    probabilities = torch.nn.functional.softmax(output, dim=1)
    confidence = probabilities.max().item() * 100

print("\nSingle Sample Test Results:")
print(f"Audio File: {audio_file}")
print(f"True Speaker ID: {true_speaker}")
print(f"Predicted Speaker ID: {predicted_speaker}")
print(f"Confidence: {confidence:.2f}%")
print(f"Prediction {'Correct' if predicted_speaker == true_speaker else 'Incorrect'}")
break


if __name__ == "__main__":
    #Run evaluation suite
    device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
    model = SpeakerCNN().to(device)
    model.load_state_dict(torch.load('best_speaker_model.pth'))

    predictions, true_labels = test_model(model, test_loader, device)
    test_random_sample(model, test_loader, device)

    specific_speaker = 17 #Test specific speaker (can be changed)
    test_file = get_audio_file_path(specific_speaker)
    if test_file:
        predicted_speaker, confidence, probabilities = predict_speaker(model, test_loader, device, test_file)
        print("\nSpecific File Test Results:")
        print(f"Test File: {test_file}")
        print(f"Predicted Speaker: {predicted_speaker}")
        print(f"Confidence: {confidence:.2f}%")
    else:
        print("Test file not found")

    #Model Performance Metrics
    print("Model Performance Metrics:")
    print(f"Test Accuracy: {accuracy:.2f}")
    print(f"Weighted Precision: {precision:.2f}")
    print(f"Weighted Recall: {recall:.2f}")
    print(f"Weighted F1-Score: {f1_score:.2f}")

```

 <ipython-input-19-dcf81d930ebe>:120: FutureWarning: You are using `torch.load` model.load_state_dict(torch.load('best_speaker_model.pth'))

Detailed Classification Report:

	precision	recall	f1-score	support
0	0.98	0.99	0.99	100
1	0.96	1.00	0.98	100
2	1.00	1.00	1.00	100
3	0.99	0.99	0.99	100
4	1.00	0.99	0.99	100
5	0.97	0.99	0.98	100
6	0.98	0.96	0.97	100
7	0.98	1.00	0.99	100
8	0.99	1.00	1.00	100
9	0.99	1.00	1.00	100
10	1.00	1.00	1.00	100
11	1.00	1.00	1.00	100
12	1.00	1.00	1.00	100
13	0.99	0.99	0.99	100
14	1.00	0.98	0.99	100
15	1.00	0.98	0.99	100
16	1.00	0.98	0.99	100
17	1.00	0.97	0.98	100
18	0.97	1.00	0.99	100
19	0.99	0.98	0.98	100
20	0.99	0.98	0.98	100
21	1.00	0.98	0.99	100
22	0.93	0.97	0.95	100
23	0.98	0.98	0.98	100
24	1.00	0.98	0.99	100
25	0.98	1.00	0.99	100
26	0.99	0.98	0.98	100
27	1.00	1.00	1.00	100
28	0.97	1.00	0.99	100
29	1.00	1.00	1.00	100
30	0.99	0.99	0.99	100
31	0.99	0.97	0.98	100
32	1.00	0.98	0.99	100
33	0.99	0.97	0.98	100
34	0.99	0.98	0.98	100
35	0.98	0.99	0.99	100
36	0.99	1.00	1.00	100
37	0.98	1.00	0.99	100
38	0.93	0.95	0.94	100
39	1.00	0.99	0.99	100
40	1.00	0.96	0.98	100
41	0.99	0.99	0.99	100
42	1.00	1.00	1.00	100
43	0.99	0.98	0.98	100
44	1.00	0.99	0.99	100
45	0.98	0.99	0.99	100
46	1.00	1.00	1.00	100

40	1.00	1.00	1.00	100
17	1.00	0.00	0.00	100